

Kill Signal Decision Card

One page. Print landscape. Run before any AI-heavy task on a legacy codebase.

Step 1 - Count Kill Signals (5 min)

For the project / module in front of you, mark each that applies:

- | # | Signal | Quick test | OK |
|---|--------------------------------------|--|-----|
| 1 | No tests (<30% coverage) | Run coverage tool; spot-check a critical service | [] |
| 2 | No documentation (tribal knowledge) | Can the senior dev be on vacation for 2 weeks and the team still ships? | [] |
| 3 | Tight coupling (>5 modules affected) | git log --stat on the last 5 PRs - any one PR touch >5 files in different domains? | [] |
| 4 | Scattered business rules | Is the AML / reconciliation / commission logic findable in one file? | [] |
| 5 | Regulatory constraints | PCI-DSS, GDPR, EU AI Act, NIS2, banking AML - at least one applies? | [] |
| 6 | Team can't evaluate output | If AI generated a crypto module right now, who could review it? | [] |
| 7 | Model-context fit (thin pretraining) | Is the codebase in a proprietary DSL, internal framework, or rare language with no public corpus? Thin model-context fit means the agent looks confident but pattern-matches the wrong corpus. | [] |
| 8 | Velocity-of-change (mid-migration) | Is the framework or a core dependency mid-migration right now? High velocity-of-change means code straddles old + new; the agent mixes idioms. | [] |

Total: ____ / 8

Step 2 - Traffic Light Action

- | Signals | Color | What you do this week |
|---------|----------|---|
| 0-1 | ? GREEN | AI-led full agentic workflow. Champions self-direct. Standard governance applies (CLAUDE.md, kill-zone). |
| 2-3 | ? YELLOW | AI-assisted. Human leads, agent supports. Every PR reviewed manually + by pr-review-toolkit. Smaller task packets (GREEN-sized only). |
| 4+ | ? RED | DO NOT use AI here yet. Invest in codebase readiness FIRST: tests, docs, decoupling. Re-classify in 6 weeks. |

Step 3 - Choose Tool by Task

- | You want to... | Use | Why |
|-------------------------------------|---|--|
| Explore an unfamiliar codebase | Claude Code | Agent loop + Read/Glob/Grep |
| Refactor Spring Boot (GREEN/YELLOW) | Claude Code + Superpowers | 6-phase methodology enforces discipline |
| Review every PR | pr-review-toolkit (6 local agents) | Specialized agents, parallel review |
| Critical-path review before merge | /ultrareview (cloud fleet, if data policy allows) | Independent verification, multi-dimensional review |

Long migration (multi-day) | Claude Code Agent View --bg | Async, 2-4 parallel sessions
Quick inline autocomplete | Cursor Tab or GitHub Copilot | Lower-friction tier
Pair-programming UI session | Cursor or Aider | UI-first ergonomics
Doc analysis (non-code) | CoWork | Same model, non-dev surface

3 Banking Examples (one per color)

- ? GREEN - Customer notification template service
- Tests at 78%; documented; isolated; no regulatory text inside the templates themselves; mid-level dev owns it; team can read the AI output.
 - Signal count: 0. AI-led. Champion can self-direct.
- ? YELLOW - Account-statement PDF generator
- Tests at 45%; logic somewhat documented; touches account-service + customer-service + format-service (3 modules); no scattered business rules; PCI-DSS scope; team can mostly evaluate.
 - Signal count: 2 (tests + regulatory). AI-assisted. Human leads; every PR through pr-review-toolkit; smaller task packets.
- ? RED - Reconciliation engine for SWIFT messages
- Tests at 12% (legacy); no docs (Andre's head); tight coupling across 11 modules; AML rules scattered in 8 files; multi-jurisdiction regulatory; only 2 people in the company can verify output.
 - Signal count: 6/8. DO NOT use AI here. First: 6-week investment in tests + docs + decoupling. Re-classify after.

Counts you'll typically see in ITSS-style projects

Project type	Typical signal count	Color
New microservice (greenfield, started 2024+) | 0-1 | GREEN
Standard Spring Boot CRUD service (2-4 years old) | 1-2 | GREEN / YELLOW boundary
Legacy monolith fragment (5+ years, partial migration) | 3-5 | YELLOW / RED boundary
Core banking transaction engine (10+ years, regulatory) | 4-7 | RED
Integration layer / ESB legacy | 4-6 | RED

When NOT to invest in readiness (RED projects)

- Sometimes the right answer is "AI never works here, the codebase is end-of-life." Markers:
- Replacement project already approved with budget
 - Codebase scheduled for sunset in <18 months
 - The 6-week readiness investment costs more than the AI savings would return in 18 months

In those cases: keep human-led, plan the sunset, redirect AI investment elsewhere.

ITSS Workshop - 28 May 2026. Apply quarterly per active project.